

Homework 4

Pan Li

CSCI-GA 3033-090: Deep Reinforcement Learning

November 21, 2020

Solution 1. Here are some descriptions of my solution for the bipedal walker problem, where I implemented an Actor-Critic-LSTM model and achieved the reward score of 240.26. I have attached a visualization of the frames. I have also attached a screenshot of the learning curve & final performance. The codes can be found in the zip file.

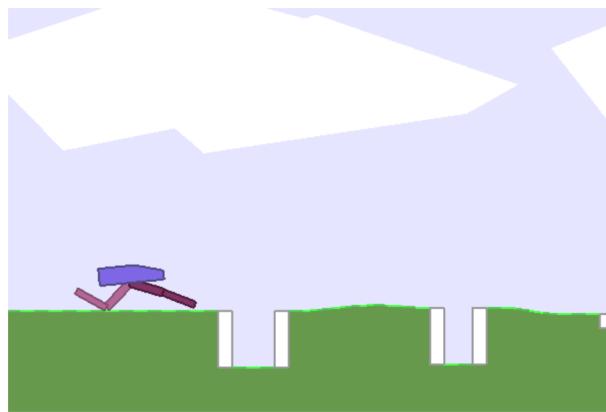


Figure 1: A Screenshot of the Frames

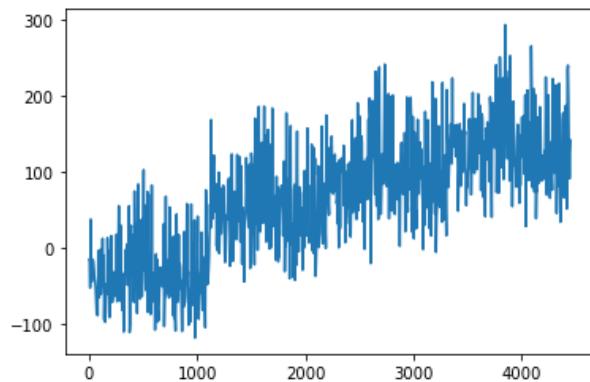


Figure 2: Plot of Reinforcement Learning Performance